

# ManiBench

## A Benchmark for Testing Visual-Logic Drift and Syntactic Hallucinations in Manim Code Generation

Oli, Nabin

February 17, 2026

### Abstract

Traditional code-generation benchmarks like HumanEval and MBPP excel at testing logic and syntax, but they fall short when code must translate into dynamic, pedagogical visuals. We introduce MANIBENCH, a specialized benchmark designed to evaluate LLM performance in generating Manim CE (Community Edition) code—a domain where temporal fidelity and version-aware API correctness are paramount. MANIBENCH addresses two critical failure modes prevalent in LLM outputs: *Syntactic Hallucinations* (generating code that is grammatically valid Python but references non-existent Manim functions, outdated or deprecated APIs, undefined classes, or calls that break under specific library versions) and *Visual-Logic Drift* (occurrences where generated visuals diverge from intended mathematical logic, such as missing events, incorrect causal relationships, timing errors, or the model struggling to animate a concept). The benchmark aims to collect 150–200 problems, launching with a pilot of 12 high-quality challenges across five difficulty levels. These span domains including calculus, linear algebra, probability, topology, and AI. Task types are uniquely structured into categories such as drift-sensitive transformations, debugging, version-conflict traps, and multi-scene narratives. To move beyond simple test-case-based checks, MANIBENCH employs a four-tier scoring framework: (1) *Executability* (Pass@1): the fraction of outputs running without exceptions or deprecated imports; (2) *Version-Conflict Error Rate*: the frequency of runs triggering mixed-API or legacy errors; (3) *Alignment Score*: the weighted fraction of required visual events that are both present and temporally accurate; and (4) *Coverage Score*: the density of pedagogical elements, including mathematical-to-visual mapping and numeric annotations. By formalizing the requirements for temporal and syntactic precision, MANIBENCH provides a foundational testbed for the next generation of automated educational content and visual-logic synthesis.

**Keywords:** Syntactic Hallucinations, Visual-Logic Drift, Manim CE, Code Generation, Benchmarking

## Contents

<b>1</b>	<b>Introduction</b>	<b>3</b>
1.1	Contributions . . . . .	3

<b>2 Problem Definition</b>	<b>3</b>
2.1 The Two Failure Modes . . . . .	3
2.1.1 Syntactic Hallucinations . . . . .	3
2.1.2 Visual-Logic Drift . . . . .	4
2.2 Evaluation Challenges . . . . .	4
<b>3 Benchmark Design</b>	<b>5</b>
3.1 Metric Definitions . . . . .	5
3.1.1 Metric 1: Executability (Pass@1) . . . . .	5
3.1.2 Metric 2: Version-Conflict Error Rate . . . . .	5
3.1.3 Metric 3: Alignment Score . . . . .	5
3.1.4 Metric 4: Coverage Score . . . . .	5
3.2 Task Categories . . . . .	6
3.3 Difficulty Levels . . . . .	6
<b>4 Benchmark Dataset</b>	<b>6</b>
4.1 Pilot Dataset: 12 Problems . . . . .	6
4.1.1 Problem 1: Colliding Blocks Compute $\pi$ . . . . .	7
4.1.2 Problem 2: Gradient Descent—How Neural Networks Learn . . . . .	7
4.1.3 Problem 3: But What Is a Convolution? . . . . .	7
4.1.4 Problem 4: Eigenvectors and Eigenvalues . . . . .	8
4.1.5 Problem 5: The Determinant . . . . .	8
4.1.6 Problem 6: The Central Limit Theorem . . . . .	8
4.1.7 Problem 7: The Medical Test Paradox (Bayes’ Theorem) . . . . .	8
4.1.8 Problem 8: Visualizing the Chain Rule . . . . .	9
4.1.9 Problem 9: Integration and the Fundamental Theorem . . . . .	9
4.1.10 Problem 10: Taylor Series . . . . .	9
4.1.11 Problem 11: The Hairy Ball Theorem . . . . .	10
4.1.12 Problem 12: The Windmill Problem . . . . .	10
<b>5 Evaluation Protocol</b>	<b>10</b>
5.1 Workflow . . . . .	10
5.2 Human Evaluation Protocol . . . . .	10
<b>6 Preliminary Results (Pilot Study)</b>	<b>11</b>
<b>7 Discussion</b>	<b>11</b>
7.1 Why ManiBench Matters . . . . .	11
7.2 Limitations and Future Work . . . . .	12
7.3 Broader Impact . . . . .	12
<b>8 Conclusion</b>	<b>12</b>
<b>A Problem Annotation Template</b>	<b>12</b>

# 1 Introduction

The rise of large language models (LLMs) has dramatically accelerated code-generation research. Benchmarks like HumanEval, MBPP, and APPS have become standard evaluation tools for assessing LLM coding ability. However, these benchmarks primarily focus on three criteria:

- **Logic correctness:** Does the code solve the algorithmic problem?
- **Syntax validity:** Does the code parse and execute without errors?
- **Output matching:** Do computed results match expected values?

These criteria are insufficient for domains where code generates continuous, time-dependent visual outputs. Manim, a Python animation engine created by Grant Sanderson (3Blue1Brown), generates mathematical animations by composing scene objects, applying transformations, and controlling timing. A Manim script can be *syntactically valid* yet produce:

- **Incorrect visual semantics:** an animation that moves in the wrong direction;
  - **Timing misalignments:** events that occur out of order or at wrong times;
  - **Pedagogical failure:** an animation that obscures rather than clarifies the concept.
- Additionally, Manim exists in two major versions:
- **Manim CE (Community Edition):** open-source, actively maintained, with a modern API;
  - **Manim GL (3B1B’s version):** the original version, using some deprecated constructs, hand-optimized for performance.

LLMs frequently mix APIs from both versions or reference functions that have been moved or renamed, producing code that fails under specific library versions.

## 1.1 Contributions

MANIBENCH makes three key contributions:

1. **Formalized Visual-Logic Metrics.** We define an *Alignment Score* and a *Coverage Score* to capture whether generated animations match pedagogical intent, beyond mere syntactic validity.
2. **Version-Aware Evaluation.** We explicitly test version-conflict errors and deprecated API usage, measuring whether code adheres to a specific Manim version’s API contract.
3. **Curated Pilot Dataset.** We provide 12 hand-crafted benchmark problems drawn from 3Blue1Brown’s published videos, with detailed visual-event specifications and annotations.

# 2 Problem Definition

## 2.1 The Two Failure Modes

### 2.1.1 Syntactic Hallucinations

An LLM generates code that:

- references non-existent classes (e.g., `VMobject` with incorrect spelling);
- uses deprecated functions (e.g., `mobject.scale()` instead of `mobject.scale_to_fit_width()`);
- calls methods with incorrect signatures;

- mixes Manim GL syntax with Manim CE (e.g., using OpenGL-specific rendering calls in CE).

**Example:**

```

1 # HALLUCINATED: class does not exist
2 circle = MCircle(color=BLUE) # Should be Circle
3
4 # HALLUCINATED: deprecated method
5 circle.apply_matrix([[1, 0], [0, 1]])
6 # CE removed in favor of apply_complex_function

```

Listing 1: Syntactic hallucination examples.

### 2.1.2 Visual-Logic Drift

An LLM generates code that:

- omits required visual events (e.g., gradient descent step without showing dot movement);
- implements events in the wrong order (e.g., loss curve updates before parameter updates);
- uses incorrect timing (animations too fast, pauses missing);
- fails to show causal relationships (e.g., showing a result without showing the derivation).

**Example:**

```

1 # DRIFTED: Gradient descent without showing step updates
2 def construct(self):
3     # Shows loss curve but dot doesn't move downhill
4     loss_curve.animate.points = new_points
5     # Missing: dot.animate.move_to(new_point)

```

Listing 2: Visual-logic drift example.

## 2.2 Evaluation Challenges

**Subjectivity of “correct.”** What counts as a correct gradient descent animation? Must the dot move along the loss surface? Must the curve update dynamically? Must the learning rate shrink?

**Version fragmentation.** A script that passes in Manim CE may fail in Manim GL. We must specify which version(s) the code targets.

**Temporal semantics.** Unlike static code output (e.g., classification accuracy), animations have temporal semantics. An event can be present but timed incorrectly, creating pedagogical failure.

## 3 Benchmark Design

### 3.1 Metric Definitions

#### 3.1.1 Metric 1: Executability (Pass@1)

**Definition 3.1** (Executability). *The fraction of generated outputs that run without raising exceptions or using deprecated imports:*

$$\text{Executability} = \frac{\text{number of successful executions}}{\text{number of total attempts}}. \quad (1)$$

**Success criteria:**

- Script completes without runtime exception.
- No deprecated imports detected (scanned via regex or AST analysis).
- No warnings from Manim's deprecation system.

**Failure cases:**

- Import error (e.g., `from manim import NonExistent`).
- Runtime `AttributeError` (e.g., `mobject.invalid_method()`).
- Type error (e.g., passing the wrong type to a function).
- Unhandled exception during scene rendering.

#### 3.1.2 Metric 2: Version-Conflict Error Rate

**Definition 3.2** (Version-Conflict Error Rate). *The frequency with which generated code triggers errors specific to version constraints:*

$$\text{VCER} = \frac{\text{number of mixed-API or legacy errors}}{\text{number of total attempts}}. \quad (2)$$

**Tracked errors:**

- GL-specific syntax in CE code.
- CE-only syntax in GL code.
- Calls to renamed or moved functions.
- Signature mismatches due to API evolution.

#### 3.1.3 Metric 3: Alignment Score

**Definition 3.3** (Alignment Score). *The weighted fraction of required visual events that are both present and temporally accurate:*

$$\text{Alignment} = \frac{\sum_i w_i \cdot p_i \cdot t_i}{\sum_i w_i}, \quad (3)$$

where  $w_i$  is the importance weight of event  $i$  ( $0 \leq w_i \leq 1$ ),  $p_i = 1$  if event  $i$  is present (0 otherwise), and  $t_i = 1$  if event  $i$  occurs at the expected time (0 otherwise).

#### 3.1.4 Metric 4: Coverage Score

**Definition 3.4** (Coverage Score). *The density of pedagogical elements (mathematical explanation, visual mapping, numeric evidence):*

$$\text{Coverage} = \frac{\text{number of elements present}}{\text{number of required elements}}. \quad (4)$$

Required elements depend on the problem domain and may include textual labels, annotations, formulas, color coding, consistent object representations, displayed numeric values, and logical scene organization.

### 3.2 Task Categories

MANIBENCH organizes problems into five categories:

1. **Direct Visualization (40%).** Prompt → Python code (classic code generation). Difficulty levels 1–3. Metric focus: Executability, Alignment Score.
2. **Drift-Sensitive (20%).** Given a script and a required temporal transformation, detect whether the visual output matches intent. Difficulty levels 2–4. Metric focus: Alignment Score, Coverage Score.
3. **Debugging (20%).** Broken code → fix (repair task). Difficulty levels 2–4. Metric focus: Executability, Alignment Score.
4. **Version-Conflict Traps (10%).** Code with tempting outdated syntax; evaluate whether the model recognizes version constraints. Difficulty levels 3–5. Metric focus: VCER, Executability.
5. **Multi-Scene Narrative (10%).** Hardest tier: multi-scene scripts combining multiple domains. Difficulty levels 4–5. Metric focus: all metrics.

### 3.3 Difficulty Levels

**Level 1 (Trivial).** Animate simple objects (circles, squares, text). Expected executability: > 95%.

**Level 2 (Basic).** Animate a transformation or a simple mathematical concept. Expected executability: 80–90%.

**Level 3 (Intermediate).** Combine multiple transformations; show a mathematical relationship. Expected executability: 70–80%.

**Level 4 (Advanced).** Multi-step derivation, temporal synchronization, pedagogical clarity. Expected executability: 50–70%.

**Level 5 (Expert).** Complex concept, multiple scenes, advanced Manim features. Expected executability: 30–50%.

## 4 Benchmark Dataset

### 4.1 Pilot Dataset: 12 Problems

The pilot dataset includes 12 hand-curated problems drawn from 3Blue1Brown’s published videos. Each problem includes:

1. a natural-language problem statement,
2. video source (YouTube link and timestamp),
3. required visual events with formal specifications,
4. difficulty level (1–5),
5. task category,
6. success criteria for Executability, Alignment, and Coverage, and
7. reference implementation notes (not shared with models; for human evaluation only).

#### 4.1.1 Problem 1: Colliding Blocks Compute $\pi$

**Metadata:** Video ID: 6dTy011fmDo; Category: Drift-Sensitive, Multi-Scene; Difficulty: 4; Domain: Physics, Numerics.

**Problem Statement.** Write Manim code to animate the collision of two blocks sliding on a frictionless surface. Block A (mass  $M$ ) starts at rest. Block B (mass  $m$ ) approaches from the left with velocity  $v_0$ . After elastic collision, count the total number of collisions. If  $m/M = 0.01$ , exactly  $\pi$  wall collisions occur. The animation must show: (1) Block A at  $x = 10$ , Block B at  $x = 0$  moving right; (2) velocity vectors above each block; (3) a collision counter incrementing at each collision; (4) velocity updates after each collision (calculated via elastic collision formulas); (5) the final state with Block B at rest and Block A moving away; and (6) text displaying the collision count.

**Required Visual Events** (weights in parentheses): blocks move and collide (0.9); collision counter increments correctly (0.8); velocity vectors update after collision (0.7); final text displays collision count (0.6).

**Success Criteria:** Executability  $\geq 0.70$ ; Alignment  $\geq 0.70$ ; Coverage  $\geq 0.75$ .

#### 4.1.2 Problem 2: Gradient Descent—How Neural Networks Learn

**Metadata:** Video ID: IHZwWFHWa-w; Category: Direct Visualization, Drift-Sensitive; Difficulty: 3; Domain: Machine Learning, Calculus.

**Problem Statement.** Create a Manim scene animating gradient descent on a 2D loss landscape. Show: (1) a parametric surface  $z = L(w_1, w_2)$ ; (2) a dot starting at a high-loss location; (3) at each step, compute  $\nabla L$ , move the dot in the direction of  $-\nabla L$ , and update a loss curve; (4) 5–10 steps of descent with diminishing step size; (5) arrows indicating gradient direction; and (6) axis labels  $w_1$ ,  $w_2$ , and “Loss.”

**Required Visual Events** (weights): surface visualized (0.8); dot at initial location (0.8); gradient arrow shown and updated (0.7); dot moves downhill (0.9); loss curve plots historical values (0.8); step size diminishes (0.6).

**Success Criteria:** Executability  $\geq 0.95$ ; Alignment  $\geq 0.75$ ; Coverage  $\geq 0.80$ .

#### 4.1.3 Problem 3: But What Is a Convolution?

**Metadata:** Video ID: KuXjwB4LzSA; Category: Direct Visualization, Drift-Sensitive; Difficulty: 3; Domain: Signal Processing, Linear Algebra.

**Problem Statement.** Animate the convolution operation between a signal and a kernel. Show: (1) a 1D signal plotted on a horizontal axis; (2) a 1D kernel displayed as a sliding window; (3) the window moving left-to-right along the signal; (4) element-wise products highlighted at each position; (5) the integral accumulating in a separate output graph; and (6) the output graph building up point-by-point.

**Required Visual Events** (weights): signal visualized (0.8); kernel visualized (0.8); window moves through signal (0.9, *critical*); product highlighted (0.7); integral accumulates (0.8); output graph builds dynamically (0.8).

**Success Criteria:** Executability  $\geq 0.90$ ; Alignment  $\geq 0.80$ ; Coverage  $\geq 0.75$ .

#### 4.1.4 Problem 4: Eigenvectors and Eigenvalues

**Metadata:** Video ID: PFDu9oVAE-g; Category: Direct Visualization; Difficulty: 4; Domain: Linear Algebra, Transformations.

**Problem Statement.** Animate how eigenvectors behave under a  $2 \times 2$  matrix transformation. Show: (1) a 2D coordinate grid with basis vectors  $\mathbf{e}_1$  and  $\mathbf{e}_2$ ; (2) a matrix  $A$  visualized as a grid deformation; (3) most vectors rotate and change length; (4) eigenvectors only change length (stay on the same line); (5) color-coded eigenvectors; (6) eigenvalues  $\lambda_1$  and  $\lambda_2$  displayed; and (7) transformation applied smoothly over two seconds.

**Required Visual Events** (weights): grid visualized (0.8); basis vectors highlighted (0.7); transformation applied (0.9); eigenvectors identified and colored (0.8); eigenvalue labels shown (0.7); eigenvectors remain collinear (0.8).

**Success Criteria:** Executability  $\geq 0.85$ ; Alignment  $\geq 0.75$ ; Coverage  $\geq 0.80$ .

#### 4.1.5 Problem 5: The Determinant

**Metadata:** Video ID: Ip3X9L0h2dk; Category: Direct Visualization; Difficulty: 2; Domain: Linear Algebra, Visualization.

**Problem Statement.** Animate the geometric interpretation of the determinant. Show: (1) a unit parallelogram defined by basis vectors; (2) a  $2 \times 2$  matrix applied to the parallelogram; (3) smooth transformation; (4) labels for original and new area; and (5) the numerical value of  $\det(A)$  updating during transformation.

**Required Visual Events** (weights): original parallelogram (0.8); matrix displayed (0.7); parallelogram transforms (0.9); new area labeled (0.8);  $\det(A)$  value displayed (0.8).

**Success Criteria:** Executability  $\geq 0.95$ ; Alignment  $\geq 0.85$ ; Coverage  $\geq 0.90$ .

#### 4.1.6 Problem 6: The Central Limit Theorem

**Metadata:** Video ID: zeJD6dqJ5lo; Category: Direct Visualization, Drift-Sensitive; Difficulty: 3; Domain: Probability, Statistics.

**Problem Statement.** Animate the Central Limit Theorem by showing how the distribution of sample means approaches a normal distribution. Show: (1) a histogram of samples from an arbitrary distribution; (2) repeated drawing of random samples with computed means; (3) a second histogram morphing from flat to bell-shaped; (4) a normal distribution overlay; and (5) explanatory text.

**Required Visual Events** (weights): original distribution (0.7); samples drawn (0.7); sample means computed and plotted (0.8); histogram of means builds (0.9); histogram converges to normal shape (0.8); normal curve overlay (0.7).

**Success Criteria:** Executability  $\geq 0.85$ ; Alignment  $\geq 0.75$ ; Coverage  $\geq 0.70$ .

#### 4.1.7 Problem 7: The Medical Test Paradox (Bayes' Theorem)

**Metadata:** Video ID: 1G4VkPoG3ko; Category: Direct Visualization; Difficulty: 2; Domain: Probability, Bayes' Theorem.

**Problem Statement.** Animate Bayes' theorem using the “Bayes box” visualization. Show: (1) a rectangle divided into four quadrants representing joint probabilities; (2) hypothetical counts populated; (3) animated division showing test-positive counts; (4) highlighted true-positive region; (5) step-by-step calculation of  $P(\text{sick} \mid +)$ ; and (6) final probability with paradox explanation.

**Required Visual Events** (weights): rectangle divided (0.8); populations labeled (0.7); populations animated (0.8); calculation shown step-by-step (0.8); final probability displayed (0.8).

**Success Criteria:** Executability  $\geq 0.95$ ; Alignment  $\geq 0.80$ ; Coverage  $\geq 0.85$ .

#### 4.1.8 Problem 8: Visualizing the Chain Rule

**Metadata:** Video ID: YG15m2VwSjA; Category: Direct Visualization; Difficulty: 3; Domain: Calculus, Function Composition.

**Problem Statement.** Animate the chain rule using function composition. Show two functions  $g(x)$  and  $f(u)$  where  $y = f(g(x))$ . Demonstrate how a small change  $dx$  propagates through  $g$  to produce  $du$ , then through  $f$  to produce  $dy$ , yielding  $\frac{d}{dx}[f(g(x))] = f'(g(x)) \cdot g'(x)$ .

**Required Visual Events** (weights): function  $g$  plotted (0.7); function  $f$  plotted (0.7); input and output labeled (0.7); small change  $dx$  shown (0.8); change propagates through  $g$  (0.8); change propagates through  $f$  (0.8); composition formula displayed (0.7).

**Success Criteria:** Executability  $\geq 0.85$ ; Alignment  $\geq 0.75$ ; Coverage  $\geq 0.75$ .

#### 4.1.9 Problem 9: Integration and the Fundamental Theorem

**Metadata:** Video ID: rfG8ce4nNh0; Category: Direct Visualization; Difficulty: 3; Domain: Calculus, Integration.

**Problem Statement.** Animate the Fundamental Theorem of Calculus. Show  $f(x)$  and  $f'(x)$ ; animate the area under  $f'(x)$  accumulating via a sweep from left to right; display a graph of the accumulated area equaling  $f(x)$ ; and demonstrate that  $\int_0^x f'(t) dt = f(x) - f(0)$ .

**Required Visual Events** (weights):  $f$  visualized (0.8);  $f'$  visualized (0.8); sweep/accumulation animated (0.9); accumulated area displayed dynamically (0.8); Fundamental Theorem formula shown (0.7).

**Success Criteria:** Executability  $\geq 0.90$ ; Alignment  $\geq 0.80$ ; Coverage  $\geq 0.80$ .

#### 4.1.10 Problem 10: Taylor Series

**Metadata:** Video ID: 3d6DsjIBzJ4; Category: Direct Visualization; Difficulty: 4; Domain: Calculus, Series.

**Problem Statement.** Animate the Taylor series expansion of  $e^x$  (or  $\sin x$ ). Plot the original function in black. Progressively add partial sums  $P_0(x)$ ,  $P_1(x)$ ,  $P_2(x)$ ,  $P_3(x)$ , ... with distinct colors. Show numerical coefficients and a convergence message. Animate 5–8 terms.

**Required Visual Events** (weights): original function plotted (0.8); partial sums  $P_0, P_1, \dots$  added progressively (0.9); each term colored and labeled (0.8); approximation improves visually (0.8); convergence demonstrated (0.8).

**Success Criteria:** Executability  $\geq 0.80$ ; Alignment  $\geq 0.75$ ; Coverage  $\geq 0.80$ .

#### 4.1.11 Problem 11: The Hairy Ball Theorem

**Metadata:** Video ID: BHdbsHFs2P0; Category: Direct Visualization; Difficulty: 5; Domain: Topology, Vector Fields.

**Problem Statement.** Animate the Hairy Ball Theorem: a continuous tangent vector field on the 2-sphere must vanish at least once. Show a 3D sphere with a tangent vector field, attempt continuous “combing,” and highlight the inevitable “bald spot.”

**Required Visual Events** (weights): sphere rendered in 3D (0.9); vector field visualized (0.9); combing attempted (0.8); discontinuity evident (0.8); bald spot highlighted (0.7).

**Success Criteria:** Executability  $\geq 0.70$ ; Alignment  $\geq 0.65$ ; Coverage  $\geq 0.60$ .

#### 4.1.12 Problem 12: The Windmill Problem

**Metadata:** Video ID: M64HUIJFTZM; Category: Drift-Sensitive, Multi-Scene; Difficulty: 4; Domain: Geometry, Combinatorics.

**Problem Statement.** Animate the windmill problem: given  $n$  points in general position, a rotating line sweeps continuously through at least two points. Show the line rotating, pivoting at geometrically correct moments, and completing a  $180^\circ$  rotation.

**Required Visual Events** (weights): points visualized (0.8); line passes through two points (0.9); line rotates (0.9); pivot events at correct times (0.8); two-point contact maintained (0.8);  $180^\circ$  rotation completed (0.7).

**Success Criteria:** Executability  $\geq 0.75$ ; Alignment  $\geq 0.70$ ; Coverage  $\geq 0.65$ .

## 5 Evaluation Protocol

### 5.1 Workflow

1. **Code Generation.** Prompt an LLM with the problem statement.
2. **Execution.** Run the generated code in Manim CE (or GL, as specified).
3. **Metric Collection.** Executability (binary); Version-Conflict Errors (binary); Alignment Score (manual or heuristic); Coverage Score (checklist-based).
4. **Aggregation.** Compute Pass@1, aggregate Alignment and Coverage.

### 5.2 Human Evaluation Protocol

For Alignment and Coverage scores, we employ structured human evaluation:

1. Watch the rendered animation.
2. Check off each required visual event as present or absent.
3. Note timing: are events synchronized correctly?
4. Assess pedagogical clarity: does the animation explain the concept?

5. Provide Alignment Score (0.0–1.0) and Coverage Score (0.0–1.0).

**Disagreement Resolution.** Two independent reviewers score each output. If disagreement exceeds 0.15, a third reviewer breaks the tie. We report inter-rater agreement via Krippendorff’s  $\alpha$  or Cohen’s  $\kappa$ .

## 6 Preliminary Results (Pilot Study)

**Setup:** Models: GPT-4o, Claude 3.5 Sonnet. Prompting: zero-shot (single problem prompt, no examples). Manim version: CE. Trials: 3 per model  $\times$  12 problems = 72 runs.

Table 1: Pilot study results across 12 benchmark problems. Executability is reported as the percentage of runs completing without error. Alignment and Coverage are averaged over successful runs; “—” indicates no successful execution.

Problem	Diff.	Exec. (GPT-4o)	Exec. (Claude)	Avg Align.	Avg Cov.
1. Colliding Blocks	4	67%	67%	0.52	0.58
2. Gradient Descent	3	100%	100%	0.73	0.82
3. Convolution	3	0%	33%	—	—
4. Eigenvectors	4	33%	33%	0.40	0.50
5. Determinant	2	100%	100%	0.85	0.88
6. CLT	3	0%	33%	—	—
7. Medical Test	2	100%	100%	0.78	0.85
8. Chain Rule	3	67%	67%	0.60	0.68
9. Integration	3	0%	33%	—	—
10. Taylor Series	4	33%	33%	0.35	0.42
11. Hairy Ball	5	0%	0%	—	—
12. Windmill	4	0%	0%	—	—

### Key Observations:

1. *Simple problems* (Determinant, Medical Test) achieve high executability (100%) and alignment ( $\geq 0.78$ ).
2. *Complex problems* (Convolution, CLT, Integration, Hairy Ball) show low executability (0–33%), suggesting API misunderstandings or missing pedagogical elements.
3. *Drift events*: Problems requiring precise timing (Gradient Descent, Windmill) show lower alignment than coverage, indicating temporal synchronization issues.
4. *Version conflicts*: No GL-specific syntax was detected in this pilot; both models avoided GL references.

## 7 Discussion

### 7.1 Why ManiBench Matters

Existing benchmarks (HumanEval, APPS) measure whether code produces correct *output*. MANIBENCH measures whether code produces correct *understanding*. This distinction is critical for educational tools, where a silent failure (wrong animation) is worse than a loud failure (runtime error).

## 7.2 Limitations and Future Work

1. **Alignment Scoring** is currently manual. Future work should explore automatic alignment detection via AST analysis or video-frame comparison.
2. **Pedagogical Validation.** We do not yet validate whether animations actually teach the concept. User studies with students could address this gap.
3. **Manim API Coverage.** As Manim evolves, the benchmark should be versioned and updated accordingly.
4. **Scalability.** Moving from 12 to 150+ problems requires annotation infrastructure and community contribution.

## 7.3 Broader Impact

MANIBENCH can be used to:

- evaluate LLM educational-content generation,
- develop better prompting strategies for animation code,
- identify systematic failure modes (e.g., models struggling with temporal synchronization), and
- drive research into improving Manim API adoption in LLMs.

## 8 Conclusion

We have introduced MANIBENCH, a specialized benchmark for evaluating Manim code generation. By formalizing metrics for syntactic correctness, version compliance, visual-logic alignment, and pedagogical coverage, MANIBENCH moves beyond simple test-case evaluation to assess whether generated animations actually communicate mathematical concepts.

The 12-problem pilot dataset demonstrates both the opportunities (simple concepts: ~80% alignment) and challenges (complex temporal reasoning: ~40% alignment) in automated animation generation. With planned expansion to 150–200 problems, MANIBENCH will serve as a foundational resource for advancing LLM-driven educational content creation.

## A Problem Annotation Template

Each problem in MANIBENCH is annotated with the following structured metadata:

**problem\_id:** unique identifier (e.g., MB-001).

**title:** descriptive title.

**video\_id:** YouTube video identifier.

**category:** task category (e.g., “drift-sensitive, multi-scene”).

**difficulty\_level:** integer 1–5.

**domain:** mathematical domain(s).

**prompt:** full natural-language problem statement.

**required\_visual\_events:** list of events, each with an identifier, description, and weight.

**coverage\_requirements:** list of required pedagogical elements.

**success\_criteria:** minimum thresholds for executability, alignment, and coverage.

**exemplar\_notes:** guidance for human evaluators.